



NEXTLEAP PM FELLOWSHIP · SPOTIFY · GROWTH TEAM

Off Repeat

Spotify perfected recommendation. Discovery still stalls — because the gap was never accuracy. It's acceptance: a reason to press play on the unfamiliar.

You have On Repeat. Meet Off Repeat.

► **Live prototype + AI workflow** · github.com/arjuncoolitr/spotify-discovery-engine

Figures: Spotify reported MAU / Premium; repeat-listening & genre-diversity from published personalization research. Reviews: 7,296 analyzed via an AI review engine.

Spotify perfected recommendation. Discovery still stalls.

~675M

monthly active users (reported)

263M

Premium subscribers

~30%

of listening is repeat / familiar

-12%

genre diversity in 6 mo · Discover Weekly (research)

THE DISCOVERY PARADOX

Spotify owns one of the world's best recommenders — so the bottleneck was never finding new music. It's that engaged listeners won't act on it.

- **Algorithm exploits familiarity** — the more you listen, the tighter the loop
- **Unfamiliar = unrewarded** — no reason to trust a new pick, so users skip
- **Loyalty backfires** — heaviest users get the stalest discovery

WHAT'S MISSING

Discover Weekly, AI DJ, SongDNA — every fix improves WHAT to recommend. None built the acceptance layer: a reason to press play on the unfamiliar.

The market gap is acceptance, not accuracy.

We didn't survey 30 people. We analyzed 7,296 reviews.

AI REVIEW ENGINE = RESEARCH AT SCALE

App Store · Play Store · YouTube · Community → one grounded pipeline.
Every claim cites a real review; no fabrication.

7,296

reviews analyzed

1,147

structured-labeled (LLM)

#1

frustration = stale / repetitive

1,933

negative (from ★ ratings)

WHAT REAL USERS SAID (CITED)

"...it locks us into a comfortable but repetitive bubble."

Community · #6077

"...stuck in a sink hole surrounded by my 2020 top songs."

Community · #6072

"Repeating the same 30 songs out of 3000 isn't shuffle."

App Store · #4896

PRIMARY + SECONDARY

Planned: 5–6 interviews with Engaged Explorers (Premium power users).
[slot in verbatim quotes]

Secondary research:

- Discover Weekly cut genre diversity ~12% in 6 mo
- LLMs/algorithms narrow taste as engagement rises
- Personalization → homogenization (popularity bias)

Users aren't failing to discover — they're refusing to accept.

Who they are — and the exact moment discovery breaks.

BEHAVIORAL SEGMENTATION

Comfort Seeker

loves the loop

Active Digger

self-discovers

Passive Repeater

skips new, no effort

Bubble-Trapped Explorer

WANTS new · can't accept

◀ Target · openness → · effort ↓



Aarav G. · 27 · Bengaluru · “Bubble-Trapped Explorer”

Premium · 4 yrs · 15 hrs/wk · 2,000+ saved songs

JTBD “Hear something genuinely new I'll love — without gambling my limited listening time on a skip.”

Trap recs feel like a rerun of his history · no reason to trust the unfamiliar · retreats to the same playlists.

THE USER JOURNEY — AND WHERE IT BREAKS

1 · Open

wants something new

Hopeful

2 · Served

familiar-leaning picks

Neutral

3 · Unfamiliar

no signal, no story

Wary

4 · Skip

nothing to trust

Retreat

5 · Familiar

loop tightens

Stuck

The failure is the silent skip at stage 4 — the unfamiliar arrives with no story, no reason to risk a listen; the loop tightens and the algorithm reads the skip as 'dislike,' narrowing further.

The gap isn't accuracy. It's acceptance.

The true problem

Recs arrive with no reason to trust them → users skip the unfamiliar and retreat to the familiar.

Value · users

Hear genuinely new music they'll love, without gambling a skip.

How we know

7,296 reviews: stale/repetitive = #1 discovery complaint; cross-source, cross-segment.

Who feels it

Engaged Explorers — high-tenure Premium power users who want new music but feel trapped.


Value · Spotify

Re-ignite discovery for the highest-LTV cohort → retention + session depth.

Why now

Every model upgrade improves accuracy & widens the acceptance gap. No one owns it yet.

OPPORTUNITY SIZING



Start with Engaged Explorers: highest LTV, feel the pain most, most articulate. Win them → expand to all Premium.

Three ways to break the bubble. Only one earns trust.

Steerable agent

INSTRUCT

You tell discovery what you want in plain language. Powerful — but Spotify already shipped Prompted Playlist / AI DJ.

Social vouching

VOUCH

Show 'liked by a friend'. Strong trust — but Spotify's social graph is thin; cold-start kills it.

Off Repeat ✓

DECOMPOSE / PERSUADE

AI introduces each new track with a personalized story (why YOU'll like it) + trust layer. Earns acceptance.

Basis	Steerable	Social	Off Repeat
User insight	3	3	5
Trust earned	2	3	5
AI feasibility	5	4	4
Business moat	2	2	5
Total	12	12	19

WHY OFF REPEAT WINS

- Keeps the user the decider — supports judgment, doesn't replace it
- Targets the real failure: the unfamiliar arriving with no reason to try it
- Stakes-adaptive: light for casual, rich for high-intent discovery
- Real moat: a personalized taste-narrative profile rivals can't copy

You have On Repeat. Meet Off Repeat.

Off Repeat is the friend who puts you on. Generative AI writes a unique, grounded reason to press play for every track — impossible to template or hand-author at scale. Recommendation picks the song; AI earns the play. Trust scales in three layers:

♥ Real friend

“Liked by Nitya” — strongest trust

🎵 AI taste-twin

“Popular with your taste”

📖 AI story

always — works with zero friends

► Live prototype + workflow

github.com/arjuncooliitr/spotify-discovery-engine



Off Repeat — for someone who loves Tame Impala + lo-fi

Crumb — “Locket”

NEW TO YOU

psych-pop · Brooklyn

the woozy Tame-Impala basslines you love — from a Brooklyn band you've never played.

♥ Liked by Nitya · your friend

Mild High Club — “Homage”

NEW TO YOU

psych-pop · Los Angeles

lush, jazzy psychedelia — that dreamy haze you love, at a Sunday tempo.

👤 Popular with your taste

Men I Trust — “Show Me How”

NEW TO YOU


indie · Montreal

the hushed lo-fi groove you keep on repeat — still way under the radar.

✦ Picked for you

One grounded engine, twice — to find the problem, then fix it.

① AI REVIEW ENGINE — HOW WE FOUND THE PROBLEM (RESEARCH AT SCALE)



▼ *the finding is the product spec*

② OFF REPEAT MVP — HOW WE ACT ON IT (LIVE, PER USER)



Shared backbone grounded generative AI · no-fabrication guardrail (RAG over real facts) · Groq LLM + sentence-transformer embeddings — the rigor that made the research trustworthy makes the feature trustworthy.

Edge cases handled — hallucination → RAG grounding · cold-start → AI story works alone · story fatigue → one glanceable line, full story on tap.

Did they discover — or just get served?

★ NORTH STAR · MEANINGFUL DISCOVERY RATE

= weekly listening on newly-adopted artists ÷ total. ~0 → **35% (M3)** → **60% (M12)**. Un-gameable — a discovery counts only when **SURFACED** (new to you) + **TRIED** (played ≥30s) + **KEPT** (saved & returned in 7d). Opening the feature and walking away is worth zero.

Tier	Metric	Target	Why it matters
Leading	Story-attach try-rate — A/B: story vs no story	+25%	the story, not the song, earns the play
Leading	New-artist play-through vs control	↑	earliest demand signal
Lagging	New-artist save-rate · 7-day return	↑ by M3	real adoption, not curiosity
Lagging	Taste breadth — distinct new artists / month	+30% M6	the bubble is actually widening
Guardrail	Skip-rate · session time	flat/down	protect core listening
Guardrail	Unaided-exploration rate	not ↓	we are not building a crutch

Rollout — P1 (M0–3): shadow-test on 1% of Engaged Explorers · P2 (M3–6): 10% + guardrail gates · P3 (M6–12): all Premium if discovery-rate ↑ and skip-rate flat.

The loop becomes a launchpad — Off Repeat doesn't chase better recs, it earns the play.

We'd rather name how it fails than pretend it won't.

CRITICAL

Acceptance ≠ the real driver

Much 'won't press play' is mood & effort, not missing trust — a story can add friction.

→ **Validate the driver in Part-2 interviews; lean-in surfaces only.**

CRITICAL

Trust is fragile

One salesy or slightly-wrong story and users tune the layer out for good.

→ **RAG grounding + no-fabrication guardrail; style variety; sample QA.**

CRITICAL

Payola & gaming

Pressure to inject promoted tracks; stream-farming the metric.

→ **Firewall from promotion; reward saved + returned, never raw plays.**

HIGH

AI becomes a crutch

Users stop exploring unaided — a new trust proxy that solves nothing.

→ **Measure unaided exploration; ease off narration as adoption grows.**

HIGH

Novelty decay

The first-try lift fades once the story feels normal.

→ **Track 30/90-day adoption retention; refresh story styles.**

HIGH

Cold-start / long tail

Thin metadata on the obscure tracks that matter most → hallucination risk.

→ **Confidence-gate; safe fallback cues; enrich long-tail data.**

Most likely failure — the story moves behavior in fewer moments than we hope. So we de-risk before building: validate the acceptance driver in interviews, ship only to lean-in moments, and design the feature to make itself unnecessary.